

# CS6301: Optimization in Machine Learning

## Lecture 8: Accelerated Gradient Descent and Practical Aspects of Gradient Descent

Rishabh Iyer

Department of Computer Science  
University of Texas, Dallas

<https://sites.google.com/view/cs-6301-optml/home>

February 10th, 2020



# Project and Assignment

- Project Deadline 1: Finalize on your Project Topics and partners:  
**February 15th 2020.**
- Projects can be done in Groups with 1-3 students per group
- You need to upload the following:
  - A Project Proposal File with a) Team members, b) Introduction and Motivation of the Project, and c) Expected Outcomes
  - A 5-7 slide summary of this for each group. You will have around 5 mins to present this on Monday (and possibly Wednesday) next week



- Summary of Results for Gradient Descent: Continuous, Smooth and Strong Convex
- Accelerated Gradient Descent and Lower Bounds
- Practical Implementational Aspects



# Summary of Results

- Lipschitz continuous functions (C). With  $\gamma = \frac{R}{B\sqrt{T}}$ , achieve an  $\epsilon$ -approximate solution in  $R^2 B^2 / \epsilon^2$  iterations



# Summary of Results

- Lipschitz continuous functions (C). With  $\gamma = \frac{R}{B\sqrt{T}}$ , achieve an  $\epsilon$ -approximate solution in  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Strongly Convex (CS) with  $\gamma_t = \mu(1+t)/2$  achieve an  $\epsilon$ -approximate solution in  $2B^2/\epsilon - 1$  iterations



# Summary of Results

- Lipschitz continuous functions (C). With  $\gamma = \frac{R}{B\sqrt{T}}$ , achieve an  $\epsilon$ -approximate solution in  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Strongly Convex (CS) with  $\gamma_t = \mu(1+t)/2$  achieve an  $\epsilon$ -approximate solution in  $2B^2/\epsilon - 1$  iterations
- Smooth Functions (S): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{R^2 L}{\epsilon}$  iterations.



# Summary of Results

- Lipschitz continuous functions (C). With  $\gamma = \frac{R}{B\sqrt{T}}$ , achieve an  $\epsilon$ -approximate solution in  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Strongly Convex (CS) with  $\gamma_t = \mu(1+t)/2$  achieve an  $\epsilon$ -approximate solution in  $2B^2/\epsilon - 1$  iterations
- Smooth Functions (S): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.



# Summary of Results

- Lipschitz continuous functions (C). With  $\gamma = \frac{R}{B\sqrt{T}}$ , achieve an  $\epsilon$ -approximate solution in  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Strongly Convex (CS) with  $\gamma_t = \mu(1+t)/2$  achieve an  $\epsilon$ -approximate solution in  $2B^2/\epsilon - 1$  iterations
- Smooth Functions (S): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.
- Concrete examples. Let  $L = B = 10, R = 1, \mu = 1$ . Then we have the:





# Summary of Results

- Lipschitz continuous functions (C). With  $\gamma = \frac{R}{B\sqrt{T}}$ , achieve an  $\epsilon$ -approximate solution in  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Strongly Convex (CS) with  $\gamma_t = \mu(1+t)/2$  achieve an  $\epsilon$ -approximate solution in  $2B^2/\epsilon - 1$  iterations
- Smooth Functions (S): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.
- Concrete examples. Let  $L = B = 10, R = 1, \mu = 1$ . Then we have the:
  - $\epsilon = 0.1$ , C: 10000, CS: 2000, S = 50, SS = 8.49 iterations



# Summary of Results

- Lipschitz continuous functions (C). With  $\gamma = \frac{R}{B\sqrt{T}}$ , achieve an  $\epsilon$ -approximate solution in  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Strongly Convex (CS) with  $\gamma_t = \mu(1+t)/2$  achieve an  $\epsilon$ -approximate solution in  $2B^2/\epsilon - 1$  iterations
- Smooth Functions (S): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.
- Concrete examples. Let  $L = B = 10, R = 1, \mu = 1$ . Then we have the:
  - $\epsilon = 0.1$ , C: 10000, CS: 2000, S = 50, SS = 8.49 iterations
  - $\epsilon = 0.01$ , C: 1000000, CS: 20000, S = 500, SS = 13.49 iterations



# Summary of Results

- Lipschitz continuous functions (C). With  $\gamma = \frac{R}{B\sqrt{T}}$ , achieve an  $\epsilon$ -approximate solution in  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Strongly Convex (CS) with  $\gamma_t = \mu(1+t)/2$  achieve an  $\epsilon$ -approximate solution in  $2B^2/\epsilon - 1$  iterations
- Smooth Functions (S): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.
- Concrete examples. Let  $L = B = 10, R = 1, \mu = 1$ . Then we have the:
  - $\epsilon = 0.1$ , C: 10000, CS: 2000, S = 50, SS = 8.49 iterations
  - $\epsilon = 0.01$ , C: 1000000, CS: 20000, S = 500, SS = 13.49 iterations
  - $\epsilon = 0.001$ , C: 100000000, CS: 200000, S = 5000, SS = 18.49 iterations

# Lower Bounds (No Proof)

- Case I: Lipschitz Continuous: Any black-box procedure will have an error of at least  $\frac{RB}{2(1+\sqrt{t})}$  (GD:  $\frac{RB}{\sqrt{T}}$ )



# Lower Bounds (No Proof)

- Case I: Lipschitz Continuous: Any black-box procedure will have an error of at least  $\frac{RB}{2(1+\sqrt{t})}$  (GD:  $\frac{RB}{\sqrt{T}}$ )
- Case II: Lipschitz Continuous + Strongly Convex: Any black-box procedure have an error of at least  $\frac{B^2}{2\mu T}$ , (GD:  $\frac{2B^2}{\mu(T+1)}$ )



# Lower Bounds (No Proof)

- Case I: Lipschitz Continuous: Any black-box procedure will have an error of at least  $\frac{RB}{2(1+\sqrt{t})}$  (GD:  $\frac{RB}{\sqrt{T}}$ )
- Case II: Lipschitz Continuous + Strongly Convex: Any black-box procedure have an error of at least  $\frac{B^2}{2\mu T}$ , (GD:  $\frac{2B^2}{\mu(T+1)}$ )
- Case III: Smooth: Any black box procedure have an error of at least  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  (GD:  $\frac{LR^2}{2T}$ )




# Lower Bounds (No Proof)

- Case I: Lipschitz Continuous: Any black-box procedure will have an error of at least  $\frac{RB}{2(1+\sqrt{t})}$  (GD:  $\frac{RB}{\sqrt{T}}$ )
- Case II: Lipschitz Continuous + Strongly Convex: Any black-box procedure have an error of at least  $\frac{B^2}{2\mu T}$ , (GD:  $\frac{2B^2}{\mu(T+1)}$ )
- Case III: Smooth: Any black box procedure have an error of at least  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  (GD:  $\frac{LR^2}{2T}$ )
- Case IV: Smooth + Strongly Convex: Define  $\kappa = \frac{L}{\mu}$ . Then Any black box procedure will have an error of at least  $\frac{\mu}{2} \left( \frac{\sqrt{\kappa}-1}{\sqrt{\kappa}+1} \right)^{2(T-1)}$  (GD:  $\frac{LR^2}{2} \left( 1 - \frac{\mu}{L} \right)^T = \frac{L}{2} \left( \frac{\kappa-1}{\kappa} \right)^T$ )



# Why can GD be slow?

- GD has suboptimal rates for smooth and smooth + strongly convex case.
- GD relies just on local gradient information
- Can we add some momentum from the progress made so far to push it faster towards the optimal?





# Attempt 1: Heavy Ball Momentum


- Recall standard gradient descent is  $x_{t+1} = x_t - \gamma_t \nabla f(x_t)$
- Idea of Momentum: Add inertia to the Ball:

$$x_{t+1} = x_t - \gamma_t \nabla f(x_t) + \beta_t (x_t - x_{t-1})$$


- Heavy Ball result: For smooth + strongly convex functions, the heavy ball algorithm converges in  $\frac{R^2}{2}(1 - \sqrt{\frac{1}{\kappa}})^T = \frac{L}{2}(\frac{\sqrt{\kappa}-1}{\sqrt{\kappa}})^T$  instead of  $\frac{R^2}{2}(1 - \frac{1}{\kappa})^T = \frac{L}{2}(\frac{\kappa-1}{\kappa})^T$  (GD convergence) iterations.
- Heavy Ball momentum not optimal for the Smooth case (though it is optimal for the strongly convex + smooth class).



# GD vs Momentum



gradient descent



heavy-ball method

# Nesterov's Accelerated Gradient Descent

- There is a gap of a factor of  $T$  for the Smooth case!  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  vs  $\frac{LR^2}{2T}$ !



# Nesterov's Accelerated Gradient Descent

- There is a gap of a factor of  $T$  for the Smooth case!  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  vs  $\frac{LR^2}{2T}$ !
- Nesterov's accelerated algorithm fixes this (we will not prove it in the class)



# Nesterov's Accelerated Gradient Descent

- There is a gap of a factor of  $T$  for the Smooth case!  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  vs  $\frac{LR^2}{2T}$ !
- Nesterov's accelerated algorithm fixes this (we will not prove it in the class)
- The algorithm follows the following procedure.



# Nesterov's Accelerated Gradient Descent

- There is a gap of a factor of  $T$  for the Smooth case!  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  vs  $\frac{LR^2}{2T}$ !
- Nesterov's accelerated algorithm fixes this (we will not prove it in the class)
- The algorithm follows the following procedure.
- Define  $\lambda_0 = 0$ ,  $\lambda_t = \frac{1 + \sqrt{1 + 4\lambda_{t-1}^2}}{2}$  and  $\beta_t = \frac{1 - \lambda_t}{\lambda_{t+1}}$ . Note  $\beta_t \leq 0$



# Nesterov's Accelerated Gradient Descent

- There is a gap of a factor of  $T$  for the Smooth case!  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  vs  $\frac{LR^2}{2T}$ !
- Nesterov's accelerated algorithm fixes this (we will not prove it in the class)
- The algorithm follows the following procedure.
- Define  $\lambda_0 = 0$ ,  $\lambda_t = \frac{1 + \sqrt{1 + 4\lambda_{t-1}^2}}{2}$  and  $\beta_t = \frac{1 - \lambda_t}{\lambda_{t+1}}$ . Note  $\beta_t \leq 0$
- Initialize  $x_1 = y_1$  as an arbitrary point



# Nesterov's Accelerated Gradient Descent

- There is a gap of a factor of  $T$  for the Smooth case!  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  vs  $\frac{LR^2}{2T}$ !
- Nesterov's accelerated algorithm fixes this (we will not prove it in the class)
- The algorithm follows the following procedure.
- Define  $\lambda_0 = 0$ ,  $\lambda_t = \frac{1 + \sqrt{1 + 4\lambda_{t-1}^2}}{2}$  and  $\beta_t = \frac{1 - \lambda_t}{\lambda_{t+1}}$ . Note  $\beta_t \leq 0$
- Initialize  $x_1 = y_1$  as an arbitrary point
- Step 1:  $y_{t+1} = x_t - \gamma_t \nabla f(x_t)$  (like normal GD:  $\gamma_t$  is learning rate, for e.g.  $1/L$ ).





# Nesterov's Accelerated Gradient Descent

- There is a gap of a factor of  $T$  for the Smooth case!  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  vs  $\frac{LR^2}{2T}$ !
- Nesterov's accelerated algorithm fixes this (we will not prove it in the class)
- The algorithm follows the following procedure.
- Define  $\lambda_0 = 0$ ,  $\lambda_t = \frac{1 + \sqrt{1 + 4\lambda_{t-1}^2}}{2}$  and  $\beta_t = \frac{1 - \lambda_t}{\lambda_{t+1}}$ . Note  $\beta_t \leq 0$
- Initialize  $x_1 = y_1$  as an arbitrary point
- Step 1:  $y_{t+1} = x_t - \gamma_t \nabla f(x_t)$  (like normal GD:  $\gamma_t$  is learning rate, for e.g.  $1/L$ ).
- Step 2:  $x_{t+1} = (1 - \beta_t)y_{t+1} + \beta_t y_t = y_{t+1} - \beta_t(y_{t+1} - y_t)$  (notice the similarity with heavy-ball?)



# Nesterov's Accelerated Gradient Descent

- There is a gap of a factor of  $T$  for the Smooth case!  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  vs  $\frac{LR^2}{2T}$ !
- Nesterov's accelerated algorithm fixes this (we will not prove it in the class)
- The algorithm follows the following procedure.
- Define  $\lambda_0 = 0$ ,  $\lambda_t = \frac{1 + \sqrt{1 + 4\lambda_{t-1}^2}}{2}$  and  $\beta_t = \frac{1 - \lambda_t}{\lambda_{t+1}}$ . Note  $\beta_t \leq 0$
- Initialize  $x_1 = y_1$  as an arbitrary point
- Step 1:  $y_{t+1} = x_t - \gamma_t \nabla f(x_t)$  (like normal GD:  $\gamma_t$  is learning rate, for e.g.  $1/L$ ).
- Step 2:  $x_{t+1} = (1 - \beta_t)y_{t+1} + \beta_t y_t = y_{t+1} - \beta_t(y_{t+1} - y_t)$  (notice the similarity with heavy-ball?)
- Theorem (Nesterov 1983): If  $f$  is convex and  $L$ -smooth,  
$$f(y_T) - f(x^*) \leq \frac{2LR^2}{T^2}$$



# Nesterov's Accelerated Gradient Descent

- There is a gap of a factor of  $T$  for the Smooth case!  $\frac{3L}{32} \frac{R^2}{(T+1)^2}$  vs  $\frac{LR^2}{2T}$ !
- Nesterov's accelerated algorithm fixes this (we will not prove it in the class)
- The algorithm follows the following procedure.
- Define  $\lambda_0 = 0$ ,  $\lambda_t = \frac{1 + \sqrt{1 + 4\lambda_{t-1}^2}}{2}$  and  $\beta_t = \frac{1 - \lambda_t}{\lambda_{t+1}}$ . Note  $\beta_t \leq 0$
- Initialize  $x_1 = y_1$  as an arbitrary point
- Step 1:  $y_{t+1} = x_t - \gamma_t \nabla f(x_t)$  (like normal GD:  $\gamma_t$  is learning rate, for e.g.  $1/L$ ).
- Step 2:  $x_{t+1} = (1 - \beta_t)y_{t+1} + \beta_t y_t = y_{t+1} - \beta_t(y_{t+1} - y_t)$  (notice the similarity with heavy-ball?)
- Theorem (Nesterov 1983): If  $f$  is convex and  $L$ -smooth,  
$$f(y_T) - f(x^*) \leq \frac{2LR^2}{T^2}$$
- Matches the lower bound upto constant factors!



# Summary of Results so Far...

- Lipschitz continuous functions (C):  $R^2 B^2 / \epsilon^2$  iterations



# Summary of Results so Far...

- Lipschitz continuous functions (C):  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Smooth (CS):  $2B^2 / \epsilon - 1$  iterations



# Summary of Results so Far...

- Lipschitz continuous functions (C):  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Smooth (CS):  $2B^2 / \epsilon - 1$  iterations
- Smooth Functions GD (SGD):  $\frac{R^2 L}{\epsilon}$  iterations.



# Summary of Results so Far...

- Lipschitz continuous functions (C):  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Smooth (CS):  $2B^2 / \epsilon - 1$  iterations
- Smooth Functions GD (SGD):  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth Functions AGD (SAGD):  $\sqrt{\frac{2LR^2}{\epsilon}}$  iterations



# Summary of Results so Far...

- Lipschitz continuous functions (C):  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Smooth (CS):  $2B^2 / \epsilon - 1$  iterations
- Smooth Functions GD (SGD):  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth Functions AGD (SAGD):  $\sqrt{\frac{2LR^2}{\epsilon}}$  iterations
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.





# Summary of Results so Far...

- Lipschitz continuous functions (C):  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Smooth (CS):  $2B^2 / \epsilon - 1$  iterations
- Smooth Functions GD (SGD):  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth Functions AGD (SAGD):  $\sqrt{\frac{2LR^2}{\epsilon}}$  iterations
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.
- Concrete examples. Let  $L = B = 10, R = 1, \mu = 1$ . Then we have the:



# Summary of Results so Far...

- Lipschitz continuous functions (C):  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Smooth (CS):  $2B^2 / \epsilon - 1$  iterations
- Smooth Functions GD (SGD):  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth Functions AGD (SAGD):  $\sqrt{\frac{2LR^2}{\epsilon}}$  iterations
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.
- Concrete examples. Let  $L = B = 10, R = 1, \mu = 1$ . Then we have the:
  - $\epsilon = 0.1$ , C: 10000, CS: 2000, SGD = 50, SAGD = 14.4, SS = 8.49 iterations



# Summary of Results so Far...

- Lipschitz continuous functions (C):  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Smooth (CS):  $2B^2 / \epsilon - 1$  iterations
- Smooth Functions GD (SGD):  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth Functions AGD (SAGD):  $\sqrt{\frac{2LR^2}{\epsilon}}$  iterations
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.
- Concrete examples. Let  $L = B = 10, R = 1, \mu = 1$ . Then we have the:
  - $\epsilon = 0.1$ , C: 10000, CS: 2000, SGD = 50, SAGD = 14.4, SS = 8.49 iterations
  - $\epsilon = 0.01$ , C: 1000000, CS: 20000, SGD = 500, SAGD: 44.72, SS = 13.49 iterations



# Summary of Results so Far...

- Lipschitz continuous functions (C):  $R^2 B^2 / \epsilon^2$  iterations
- Lipschitz continuous functions + Smooth (CS):  $2B^2 / \epsilon - 1$  iterations
- Smooth Functions GD (SGD):  $\frac{R^2 L}{\epsilon}$  iterations.
- Smooth Functions AGD (SAGD):  $\sqrt{\frac{2LR^2}{\epsilon}}$  iterations
- Smooth + Strongly Convex (SS): With  $\gamma = 1/L$ , achieve an  $\epsilon$ -approximate solution in  $\frac{L}{\mu} \log(\frac{R^2 L}{2\epsilon})$  iterations.
- Concrete examples. Let  $L = B = 10, R = 1, \mu = 1$ . Then we have the:
  - $\epsilon = 0.1$ , C: 10000, CS: 2000, SGD = 50, SAGD = 14.4, SS = 8.49 iterations
  - $\epsilon = 0.01$ , C: 1000000, CS: 20000, SGD = 500, SAGD: 44.72, SS = 13.49 iterations
  - $\epsilon = 0.001$ , C: 100000000, CS: 200000, SGD = 5000, SAGD = 141.42, SS = 18.49 iterations



# Does this Matter in Practice?

- Convergence results and Lower bounds are often worst case!



# Does this Matter in Practice?

- Convergence results and Lower bounds are often worst case!
- Though there exists family of functions where the bounds are tight, it is not necessary that the same intuition carries over in practice!



# Does this Matter in Practice?

- Convergence results and Lower bounds are often worst case!
- Though there exists family of functions where the bounds are tight, it is not necessary that the same intuition carries over in practice!
- Difference between Theory and Practice and the need to connect the two!



# Does this Matter in Practice?

- Convergence results and Lower bounds are often worst case!
- Though there exists family of functions where the bounds are tight, it is not necessary that the same intuition carries over in practice!
- Difference between Theory and Practice and the need to connect the two!
- Next, we will implement some of these algorithms for various ML Loss Functions!





# Gradient Descent in Practice: Basic Version

- Credits to Mark Schmidt from UBC for this (I converted his Matlab based tutorial to python)



# Gradient Descent in Practice: Basic Version

- Credits to Mark Schmidt from UBC for this (I converted his Matlab based tutorial to python)
- Let us first initialize the Logistic Loss on a dataset



# Gradient Descent in Practice: Basic Version

- Credits to Mark Schmidt from UBC for this (I converted his Matlab based tutorial to python)
- Let us first initialize the Logistic Loss on a dataset
- Next, implement a simple gradient descent algorithm.

```
def gd( funObj , w , maxEvals , alpha , ...  
X , y , lam , verbosity , freq ) :  
    ...
```

[python]



# Gradient Descent in Practice: Basic Version

- Credits to Mark Schmidt from UBC for this (I converted his Matlab based tutorial to python)
- Let us first initialize the Logistic Loss on a dataset
- Next, implement a simple gradient descent algorithm.

```
def gd( funObj , w , maxEvals , alpha , ...  
X , y , lam , verbosity , freq ) :  
    ...
```

[python]

- 'funObj' is the



# Gradient Descent in Practice: Basic Version

```
def gd(funObj,w,maxEvals,alpha,X,y,lam,verbosity):  
    [f,g] = funObj(w,X,y,lam)  
    funEvals = 1  
    funVals = []  
    while(1):  
        [f,g] = funObj(w,X,y,lam)  
        optCond = LA.norm(g, np.inf)  
        if (verbosity > 0):  
            print(funEvals,alpha,f,optCond)  
        w = w - alpha*g  
        funEvals = funEvals+1  
        if ((optCond < 1e-2) and (funEvals > maxEvals)):  
            break  
        funVals.append(f)  
    return funVals
```



# Gradient Descent in Practice: Basic Version

- Run this by invoking:

```
funV = gd(LogisticLoss ,w,200 ,1e-1,X,y,1,1,10)
```

- Try running this with different values of learning rates:  
 $\alpha = 1e-1, 1e-3, 1e-5, \dots$
- How do we find the optimal learning rate every time?
- Can there be better strategies to adapt the learning rates?
- Next, we shall see a few line search based strategies.



# Armijo Backtracking Line-Search V1

- We don't want to tune  $\alpha$  every time
- This is the idea behind line search
- Simple Line search strategy:
  - Start with a large value of  $\alpha$
  - Divide  $\alpha$  by 1/2 if it doesn't satisfy Armijo's condition:

$$f(w - \alpha g) \leq f(w) - \gamma \alpha \|g\|^2$$

- Basically find  $\alpha$  such that there is a reduction in function value by atleast  $\gamma \alpha \|g\|^2$
- Idea: Choose  $\alpha$  and  $\gamma$  such that this happens.



# Armijo Backtracking Line-Search V2

- Danger with the simple backtracking is that  $\alpha$  may quickly become very small quickly
- Easy fix: Reset  $\alpha$  every time!
- Issue with this: Too many function evaluations lost in repeated backtracking!





# Armijo Backtracking Line-Search V3

- Just halving the step size ignores the information collected during line search!
- Reduce the number of backtracks using a polynomial interpolation!
- Minimize a quadratic passing through  $f(w)$ ,  $f'(w)$  and  $f(w - \alpha g)$
- Choose  $\alpha$  using a polynomial interpolation as follows:

$$\alpha = \frac{\alpha^2 g^T g}{2(f_{curr} + \alpha g^T g - f)}$$

- Here  $f_{curr}$  is the function evaluation with the current value of  $\alpha$  and  $f$  is the function value before starting backtracking!



# Armijo Backtracking Line-Search V4

- Final Issue to fix is better initialization of  $\alpha$ .
- Initializing  $\alpha = 1$  is too large in practice
- Wasted backtracks because of this.
- Use a heuristic like  $\alpha = 1/\|g\|$
- On subsequent iterations again use a polynomial interpolation:

$$\alpha = \min(1, 2(f_{old} - f)/g^T g)$$

- A lot of this is tried empirically and based on empirical knowledge..



# Finally: Accelerated Gradient Descent

- Algorithm:

- Define  $\lambda_0 = 0$ ,  $\lambda_t = \frac{1 + \sqrt{1 + 4\lambda_{t-1}^2}}{2}$  and  $\gamma_t = \frac{1 - \lambda_t}{\lambda_{t+1}}$ .
- Note  $\gamma_t \leq 0$
- Initialize  $x_1 = y_1$  as an arbitrary point
- Step 1:  $y_{t+1} = x_t - \alpha \nabla f(x_t)$  (like normal GD)
- Step 2:  $x_{t+1} = (1 - \gamma_t)y_{t+1} + \gamma_t y_t = y_{t+1} - \gamma_t(y_{t+1} - y_t)$  (slide a little bit further than  $y_{t+1}$  towards the previous point  $y_t$ !)
- In practice, we club this with Armijo line search for tuning the learning rate  $\alpha$ .

